



## John Deardurff

Microsoft Customer Engineer (Global Technical Team)

Microsoft Certified Trainer (Regional Lead)

MVP: Data Platform (2016 – 2018)

Email: [John.Deardurff@Microsoft.com](mailto:John.Deardurff@Microsoft.com)

Twitter: [@SQLMCT](https://twitter.com/SQLMCT)

Website: [www.SQLMCT.com](http://www.SQLMCT.com)

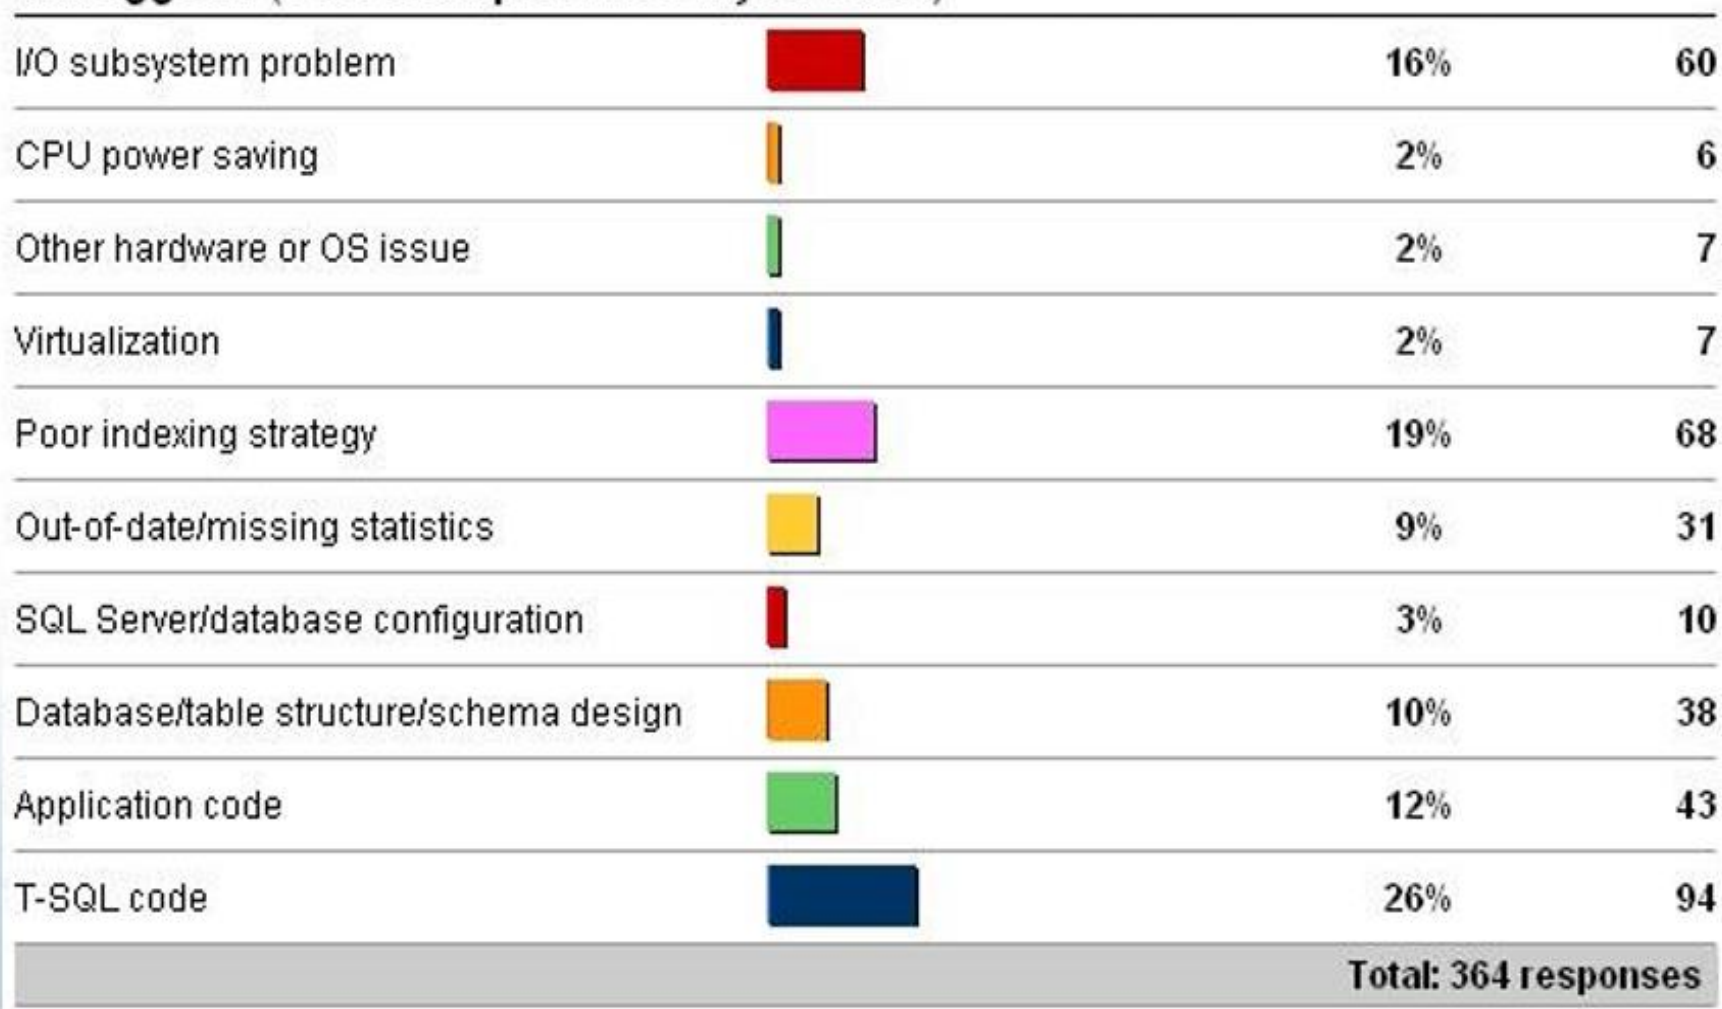
GitHub: [github.com/SQLMCT](https://github.com/SQLMCT)



# Common Cause of Performance Problems


<http://www.sqlskills.com/blogs/paul/survey-results-common-causes-of-performance-problems/>

**What were the root causes of the last few SQL Server performance problems you debugged? (Vote multiple times if you want!)**



# SQL Server Performance Killers

- Poor Indexing
- Inaccurate Statistics
- Poor Query Design
- Poor Execution Plans
- Excessive blocking and deadlocks
- Non set-based operations
- Poor database design
- Excessive fragmentation
- Non-reusable execution plans
- Frequent recompilation of queries
- Improper use of cursors
- Improper configuration of database log
- Excessive use or improper configuration of tempdb



# Two Main Functions of SQLOS

## Management

- Memory Manager
- Process Scheduler
- Synchronization
- I/O
- Support for Non-Uniform Memory Access (NUMA) and Resource Governor

## Monitoring

- Resource Monitor
- Deadlock Monitor
- Scheduler Monitor
- Lazy Writer (Buffer Pool management)
- Dynamic Management Views (DMVs)
- Extended Events
- Dedicated Administrator Connection (DAC)

# Dynamic Management Views and Functions


Category	Description
sys.dm_exec_%	Execution and connection information
sys.dm_os_%	Operating system related information
sys.dm_tran_%	Transaction management information
sys.dm_io_%	I/O related information
sys.dm_db_%	Database information

# Using Dynamic Management Objects (DMOs)

- Must reference using the sys schema
- Two basic types:
  - Real-time state information
  - Historical information

```
SELECT cpu_count, hyperthread_ratio,  
       scheduler_count, scheduler_total_count,  
       affinity_type, affinity_type_desc,  
       softnuma_configuration, softnuma_configuration_desc,  
       socket_count, cores_per_socket, numa_node_count,  
       sql_memory_model, sql_memory_model_desc  
FROM sys.dm_os_sys_info
```

# Yielding





# Thread States and Queues

**Runnable:** The thread is currently in the Runnable Queue waiting to execute. (First In, First Out).

**Running:** One active thread executing on a processor.

**Suspended:** Placed on a Waiter List waiting for a resource other than a processor. (No specific order).



# Waiting Tasks DMV

```
SELECT w.session_id, w.wait_duration_ms, w.wait_type,  
       w.blocking_session_id, w.resource_description,  
       s.program_name, t.text, t.dbid, s.cpu_time, s.memory_usage  
FROM sys.dm_os_waiting_tasks as w  
     INNER JOIN sys.dm_exec_sessions as s  
       ON w.session_id = s.session_id  
     INNER JOIN sys.dm_exec_requests as r  
       ON s.session_id = r.session_id  
     OUTER APPLY sys.dm_exec_sql_text (r.sql_handle) as t  
WHERE s.is_user_process = 1;
```

session_id	wait_duration_ms	wait_type	blocking_session_id	resource_description
58	8563	LCK_M_S	62	keylock hobtid=72057594047365120 dbid=5 id=lock1...


# Troubleshooting Wait Types

Aaron Bertrand – Top Wait Types

<https://sqlperformance.com/2018/10/sql-performance/top-wait-stats>

Paul Randal – SQL Skills Wait Types Library

<https://www.sqlskills.com/help/waits/>



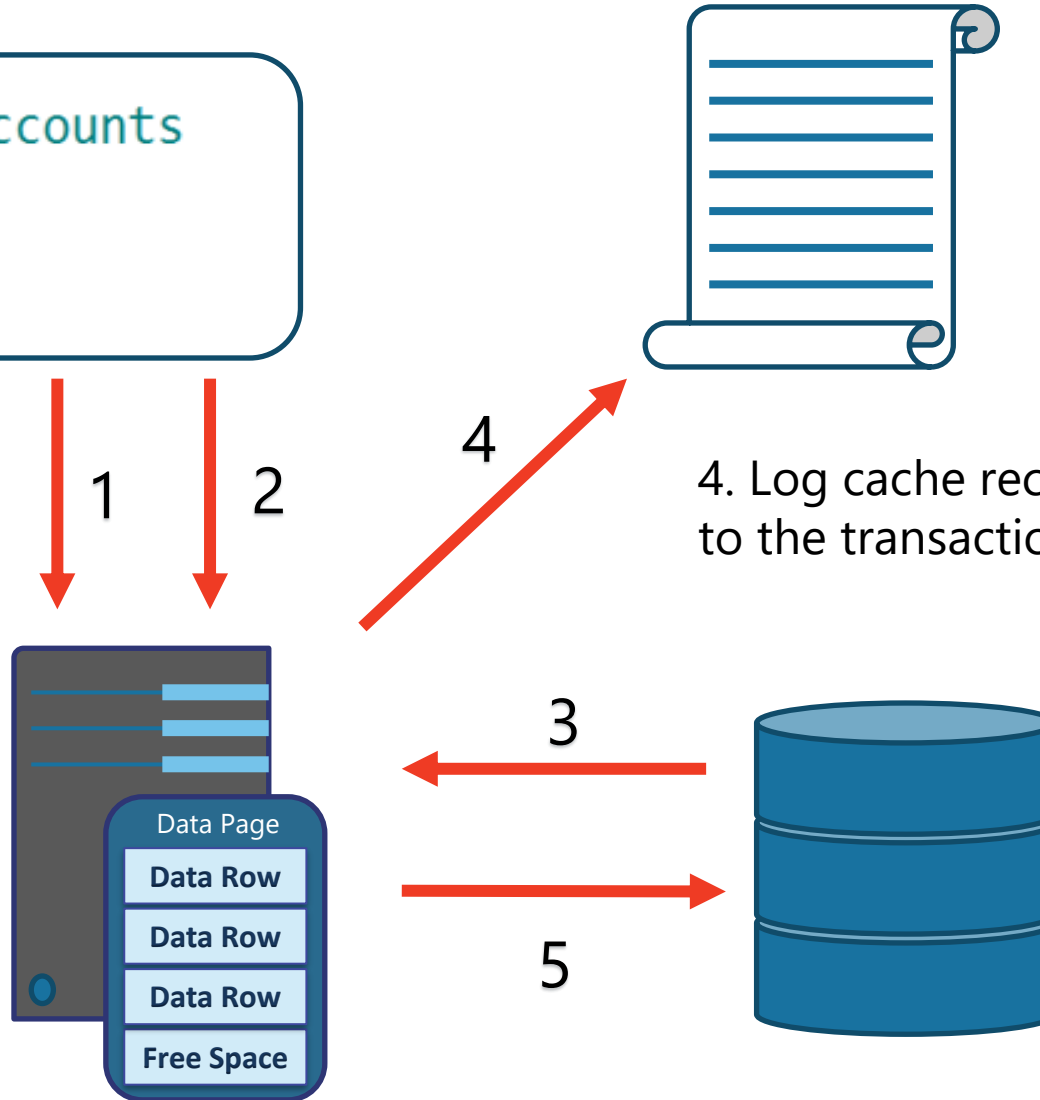
# SQL Server Disk I/O (Write-Ahead Logging)

```
UPDATE Accounting.BankAccounts  
SET Balance -= 200  
WHERE AcctID = 1
```

1. Data modification is sent to buffer cache in memory.

2. Modification is recorded in the log cache.

3. Data pages are located or read into the buffer cache and then modified.



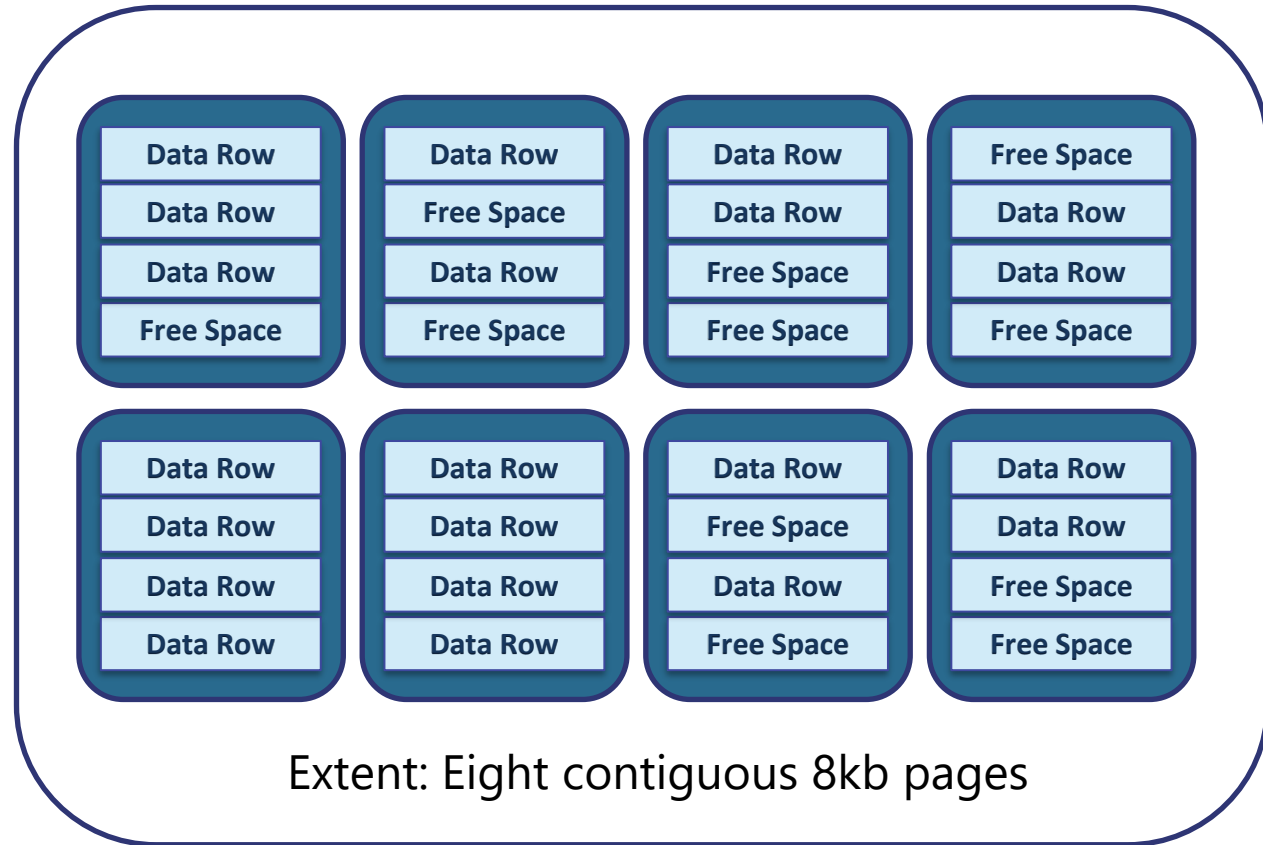
4. Log cache record is flushed to the transaction log

5. At checkpoint, dirty data pages are written to the database file.

# SQL Server Object Allocation



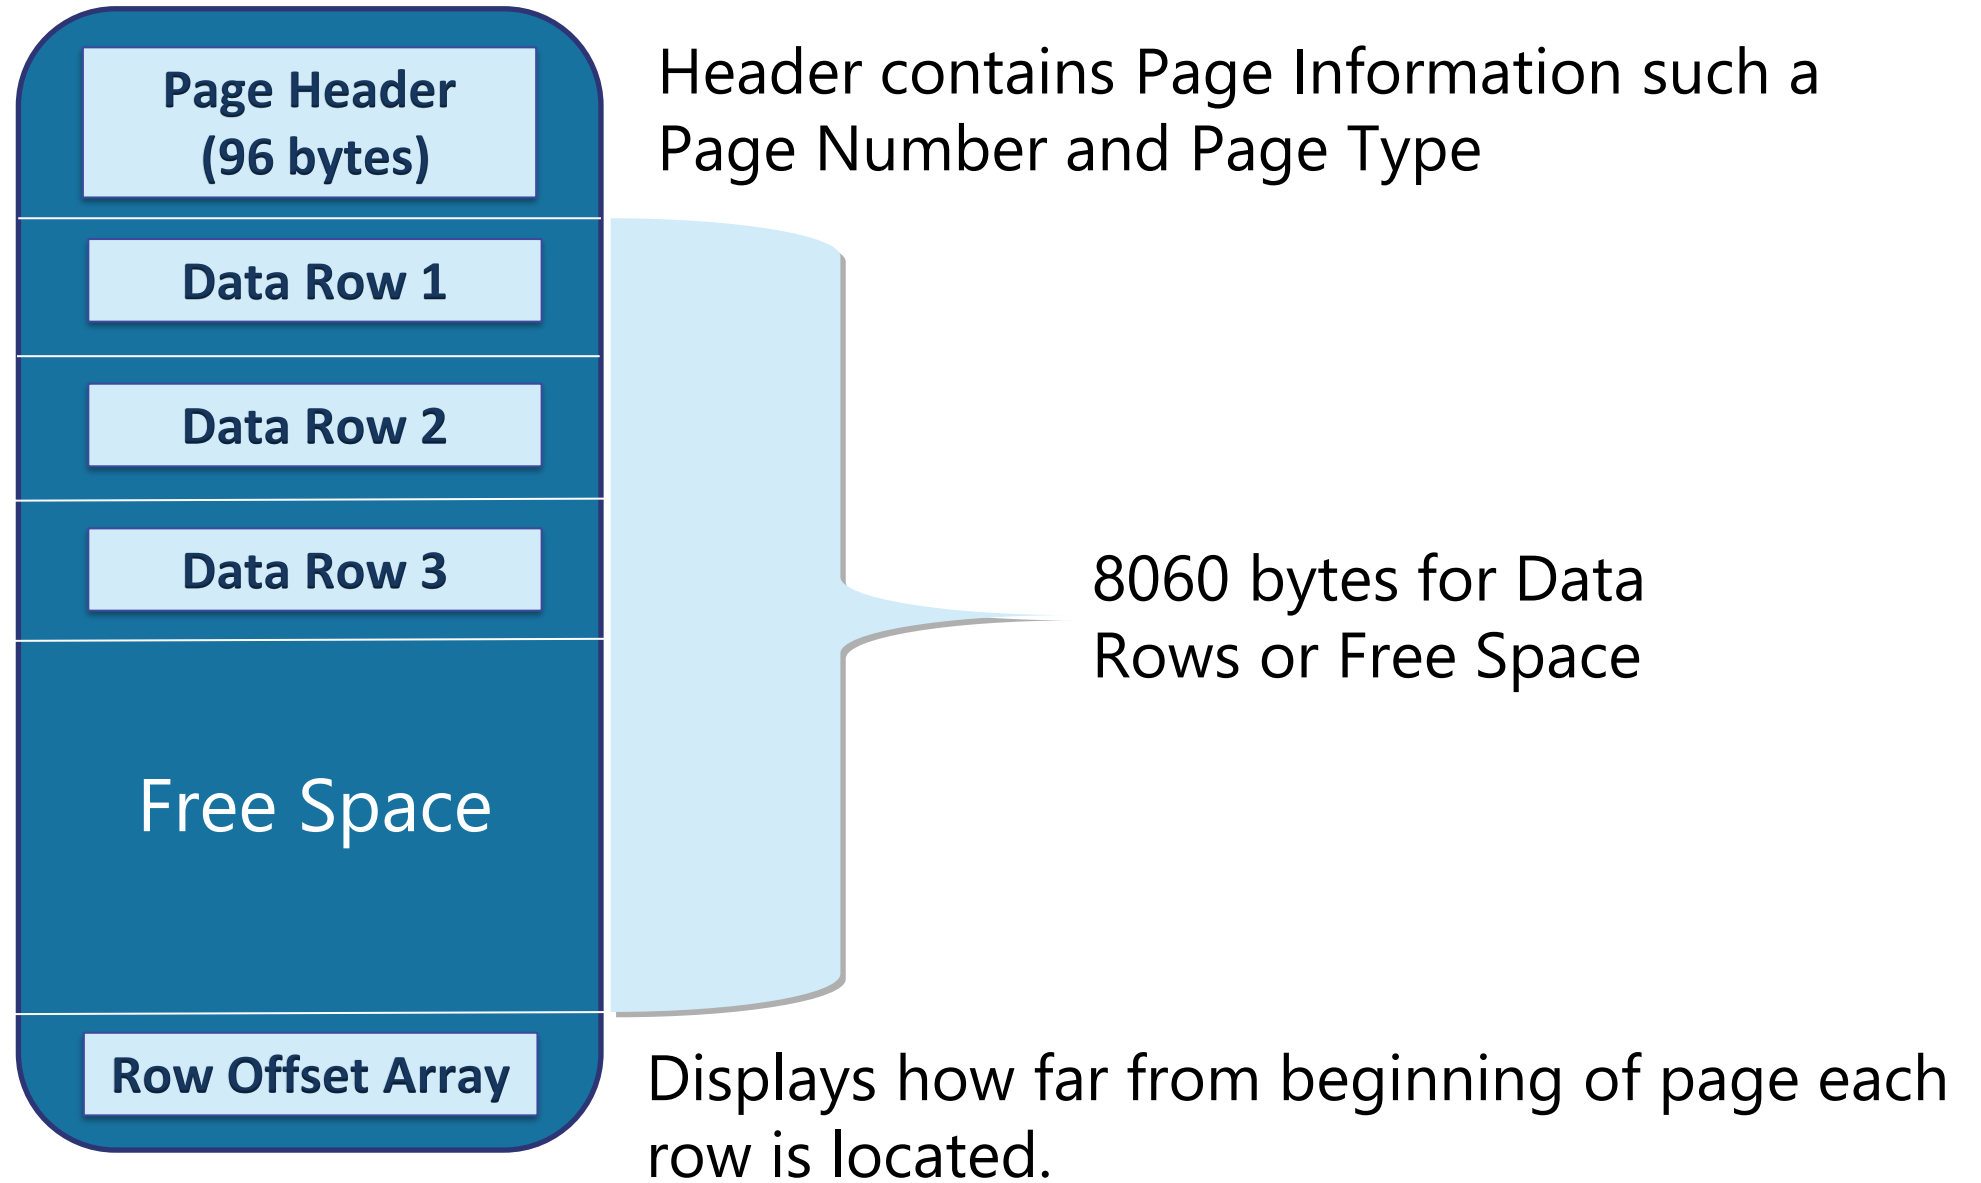
Primary Data File (.mdf)  
Secondary Data File (.ndf)



Uniform extents: Pages used by a single object.

Mixed extents: Pages used by different objects.

# Basic Page Structure



# Allocation Units

## IN\_ROW\_DATA

- Fixed length data must be store here.
- Rows cannot extend beyond pages
- Data Page is 8060 bytes

## LOB\_DATA (For out of row storage)

- varchar(max) / nvarchar(max) / varbinary(max)
- 16-byte point to out of row tree
- Uses text page to store a stream of data

## ROW\_OVERFLOW\_DATA (SLOB)

- varchar(8000) / nvarchar(4000) / varbinary(8000)
- When a column can't fit onto a page
- No control over which column overflows

# Allocation Structures

Page Types
Data (1)
Index (2)
Text (3 and 4)
Boot (13)
File Header (15)
PFS (11)
GAM (8)
SGAM (9)
IAM (10)
DIFF_MAP(16)
ML_MAP (17)

## Page Free Space (PFS)

- Tracks free space on a page (1 Byte/Page )
- Covers 64 megabytes (MB) worth of pages

## Global Allocation Map (GAM)

- Tracks which extents have been allocated (1 Bit)
- Covers 64,000 extents (4 gigabytes (GB) worth of data)

## Shared Global Allocation Map (SGAM)

- Tracks which extents are used for mixed extent allocations (1 Bit)
- Covers 64,000 extents (4 GB worth of data)

## Index Allocation Map (IAM)

- Tracks which extents are allocated to an allocation unit
- Covers 4 GB worth of data
- One IAM chain per table, per index, per partition, per allocation unit type



# Index Allocation Map Pages


SGAM:  
00011000000000

GAM:  
010110100000

Unallocated  
Extent

Mixed Extent  
with free  
pages

Unallocated  
Extent



Full  
Extent

Full  
Extent

Full  
Extent

Full  
Extent

# DBCC IND

```
USE AdventureWorks2012
DBCC TRACEON(3604) -- Print to results pane
DBCC IND(0, 'HumanResources.Employee', -1)
-- Parameter 1: Is the DatabaseName, 0 is current database
-- Parameter 2: The table name
-- Parameter 3: Index ID, -1 Shows all indexes, -2 shows only IAM Pages
```

	PageFID	PagePID	IAMFID	IAMPID	ObjectID	IndexID	PartitionNumber	PartitionID	iam_chain_type	PageType	IndexLevel	NextPageFID	NextPagePID	PrevPageFID	PrevPagePID
1	1	874	NULL	NULL	1237579447	1	1	72057594045136896	In-row data	10	NULL	0	0	0	0
2	1	875	1	874	1237579447	1	1	72057594045136896	In-row data	2	1	0	0	0	0
3	1	1048	1	874	1237579447	1	1	72057594045136896	In-row data	1	0	1	1049	0	0
4	1	1049	1	874	1237579447	1	1	72057594045136896	In-row data	1	0	1	1050	1	1048
5	1	1050	1	874	1237579447	1	1	72057594045136896	In-row data	1	0	1	1051	1	1049
6	1	1051	1	874	1237579447	1	1	72057594045136896	In-row data	1	0	1	1052	1	1050
7	1	1052	1	874	1237579447	1	1	72057594045136896	In-row data	1	0	1	1053	1	1051
8	1	1053	1	874	1237579447	1	1	72057594045136896	In-row data	1	0	1	1054	1	1052
9	1	1054	1	874	1237579447	1	1	72057594045136896	In-row data	1	0	0	0	1	1053
10	1	9287	NULL	NULL	1237579447	2	1	72057594050510848	In-row data	10	NULL	0	0	0	0
11	1	9286	1	9287	1237579447	2	1	72057594050510848	In-row data	2	0	0	0	0	0
12	1	9289	NULL	NULL	1237579447	3	1	72057594050576384	In-row data	10	NULL	0	0	0	0

Query executed successfully.	STUDENTSERVER (12.0 RTM)	STUDENTSERVER\Student ...	AdventureWorks2012	00:00:0
------------------------------	--------------------------	---------------------------	--------------------	---------

# DBCC PAGE

```
DBCC TRACEON(3604) -- Print to results pane
DBCC PAGE (0,1,0,3)
-- Parameter 1: Is the DatabaseName, 0 is current database
-- Parameter 2: The File ID
-- Parameter 3: The Page ID
-- Parameter 4: The print option, 3 is verbose
```

.00 % <

Messages

PAGE HEADER:

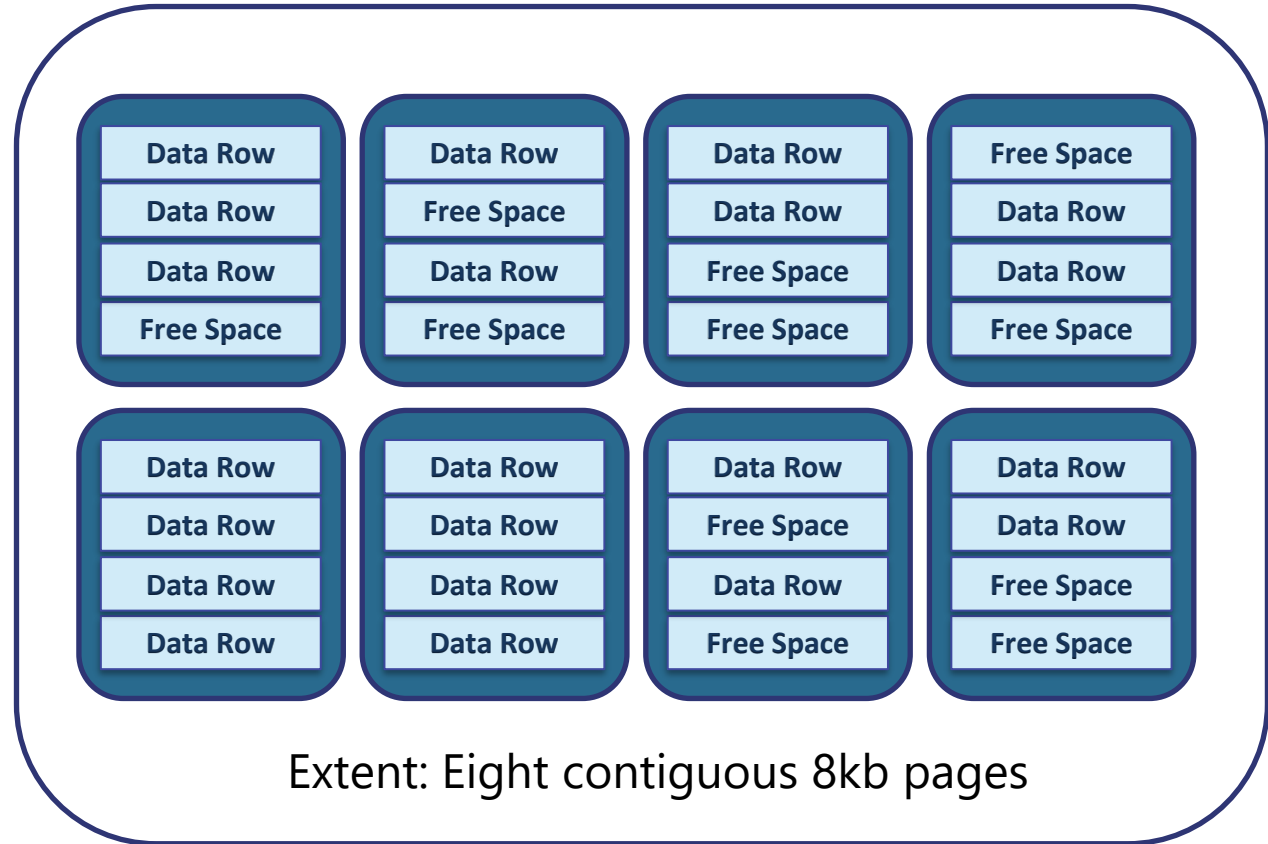
Page @0x000000027757A000

m_pageId = (1:0)	m_headerVersion = 1	m_type = 15
m_typeFlagBits = 0x0	m_level = 0	m_flagBits = 0x208
m_objId (AllocUnitId.idObj) = 99	m_indexId (AllocUnitId.idInd) = 0	Metadata: AllocUnitId = 6488064
Metadata: PartitionId = 0	Metadata: IndexId = 0	Metadata: ObjectId = 99
m_prevPage = (0:0)	m_nextPage = (0:0)	pminlen = 0
m_slotCnt = 1	m_freeCnt = 6989	m_freeData = 7831
m_reservedCnt = 0	m_lsn = (181:50952:34)	m_xactReserved = 0
m_xdesId = (0:0)	m_ghostRecCnt = 0	m_tornBits = -820886669
DB Frag ID = 1		

# How Data is stored in a Database



Primary Data File (.mdf)  
Secondary Data File (.ndf)

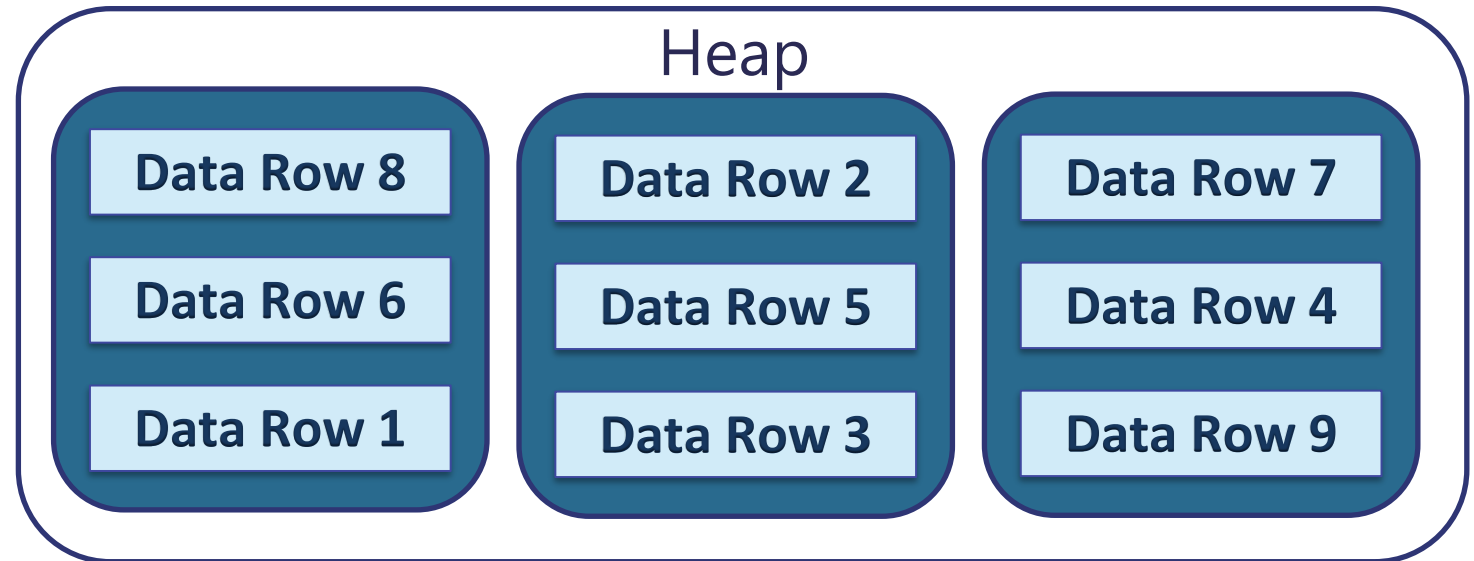


Uniform extents: Pages used by a single object.

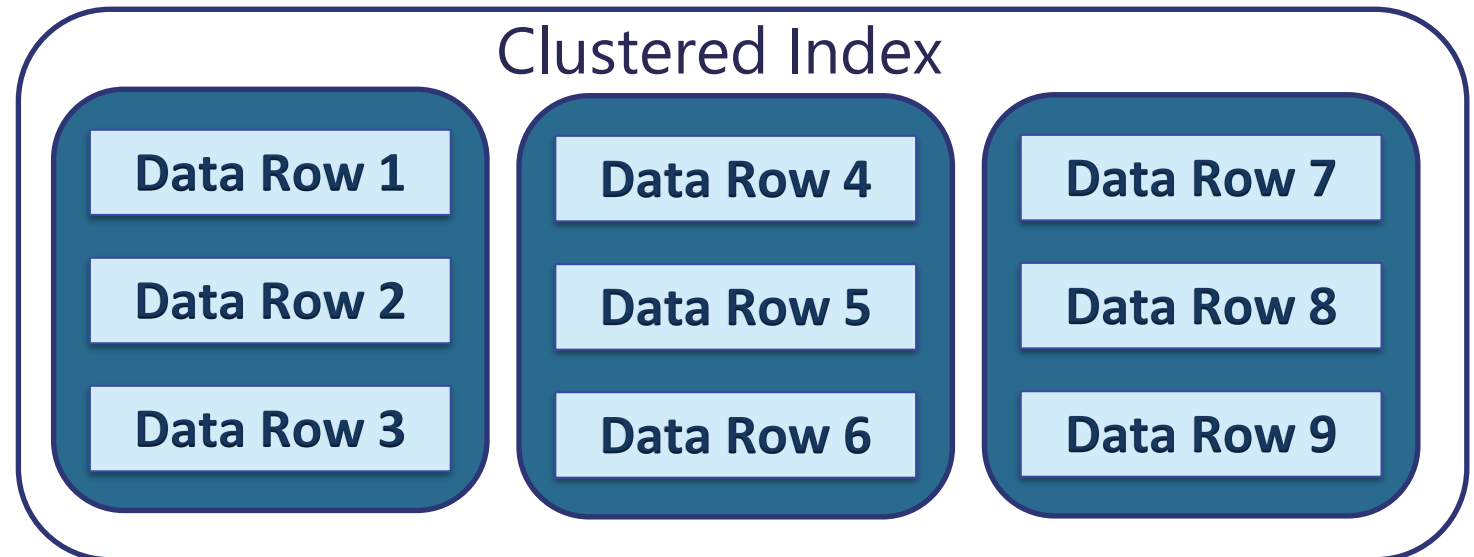
Mixed extents: Pages used by different objects.

# How Data is Stored in Data Pages

Data stored in a Heap is not stored in any order and normally does not have a Primary Key.



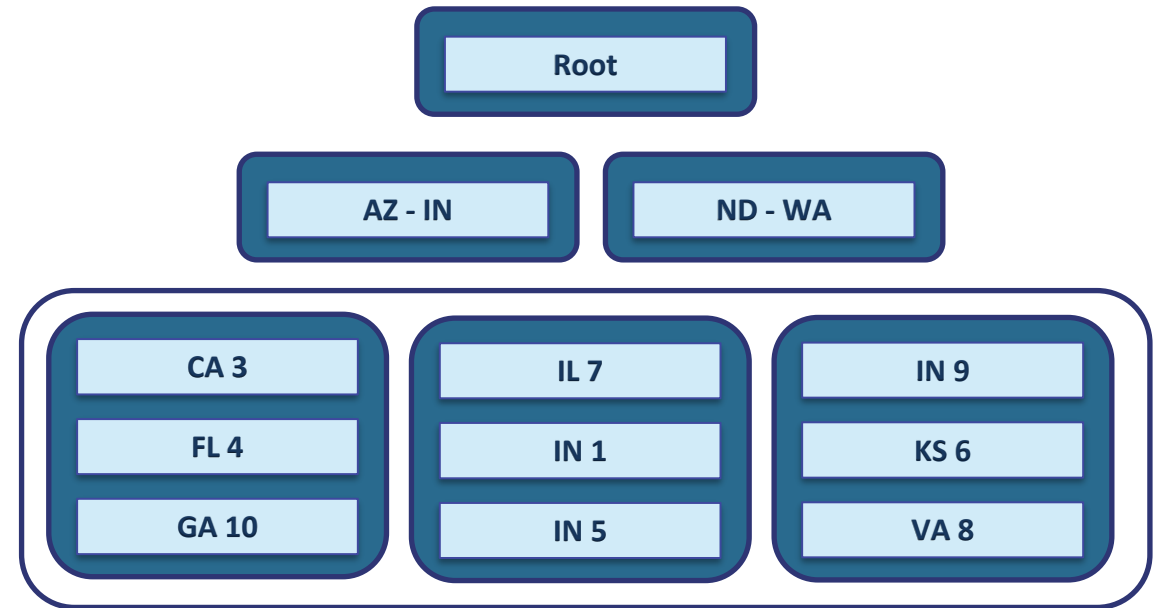
Clustered Index data is stored in sorted order by the Clustering key. In many cases, this is the same value as the Primary Key.



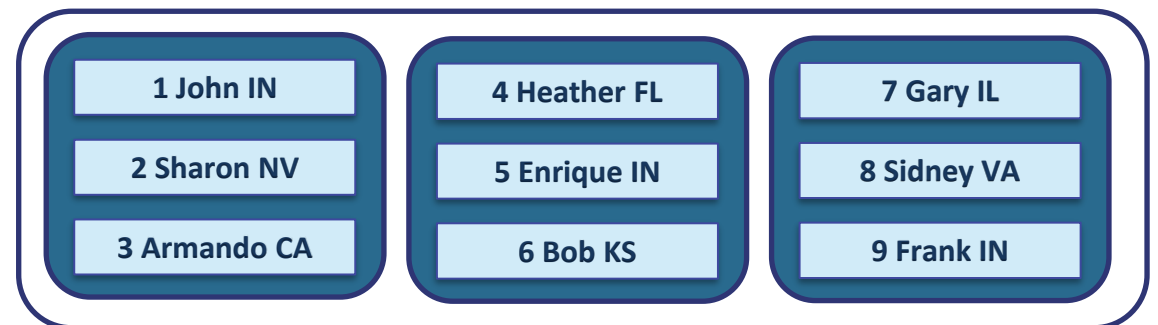
# Key or Rid Lookup

A Non-Clustered Index is built separate from the table.

The leaf level will have RID or Key values to lookup additional columns.



The data is stored in either a Clustered Index (sorted) or Heap (unsorted) table structure



# Characteristics of a Good Clustering Key

## Narrow

- Use a data type with a small number of bytes to conserve space in tables and indexes

## Unique

- To avoid SQL adding a 4-byte uniquifier

## Static

- Allows data to stay constant without constant changes which could lead to page splits

## Increasing

- Allows better write performance and reduces fragmentation issues



# How to determine Thread Stack Memory


Maximum Worker Threads  
 $512 + (\text{Processors} - 4) * 16$

\*

2mb per thread

Cores	Threads	Memory (MB)
4	512	1,024
8	576	1,152
16	704	1,408
32	960	1,920
64	1,472	2,944
80	1,728	3,456

# Dynamic Memory Management

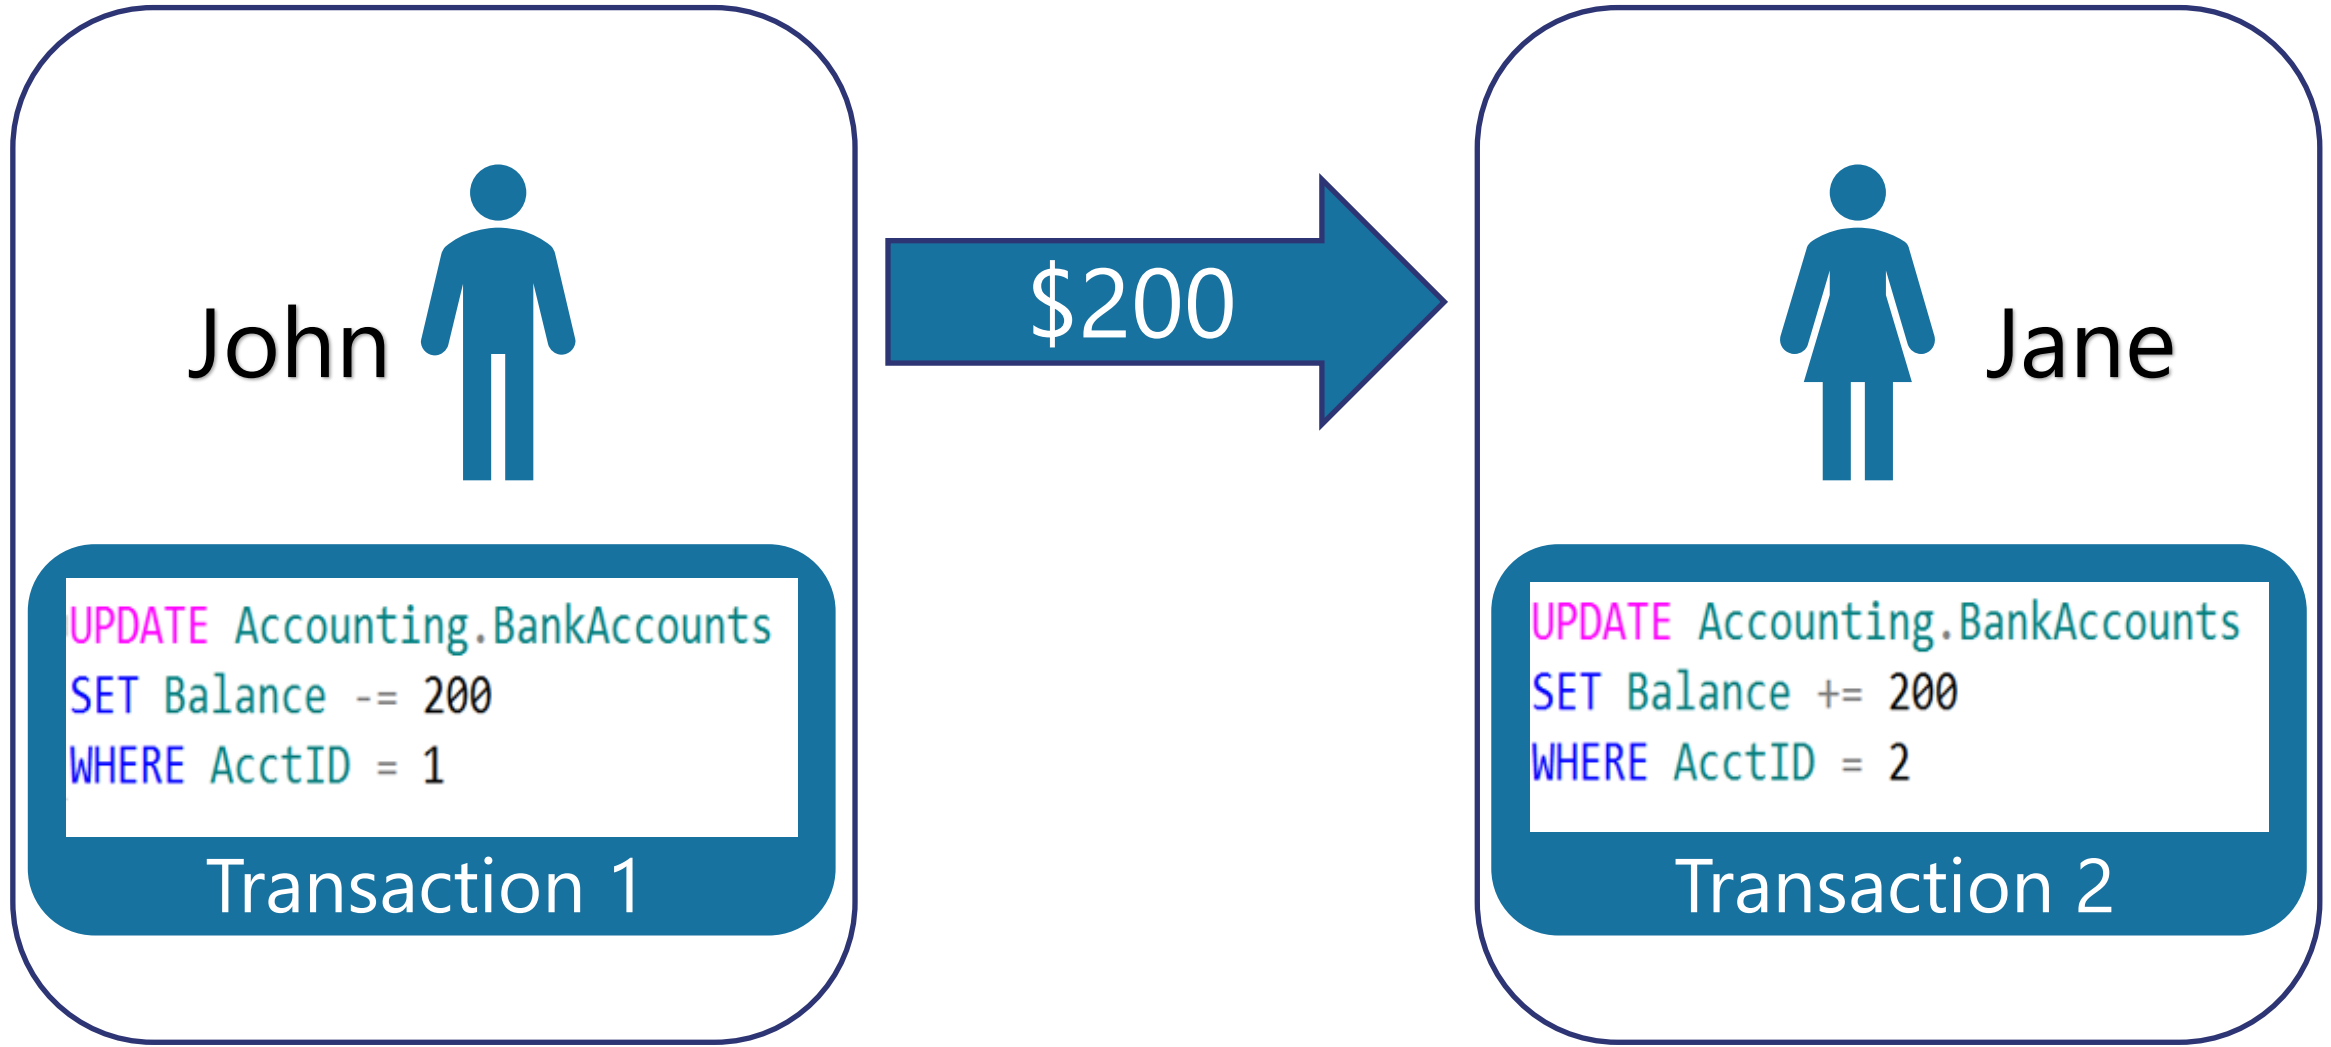


# What is a Transaction?

A transaction is a series of one or more statements that need to operate as a single logical unit of work.


To qualify as a transaction, the logical unit of work must possess all four of the ACID properties.

# Logical Units of Work – Auto Commit Transactions



# Single Logical Unit of Work – Explicit Transactions

John





```
Begin Transaction BankUpdate  
UPDATE Accounting.BankAccounts  
SET Balance -= 2/0  
WHERE AcctID = 1  
  
UPDATE Accounting.BankAccounts  
SET Balance += 200  
WHERE AcctID = 2  
Commit Transaction
```



Jane

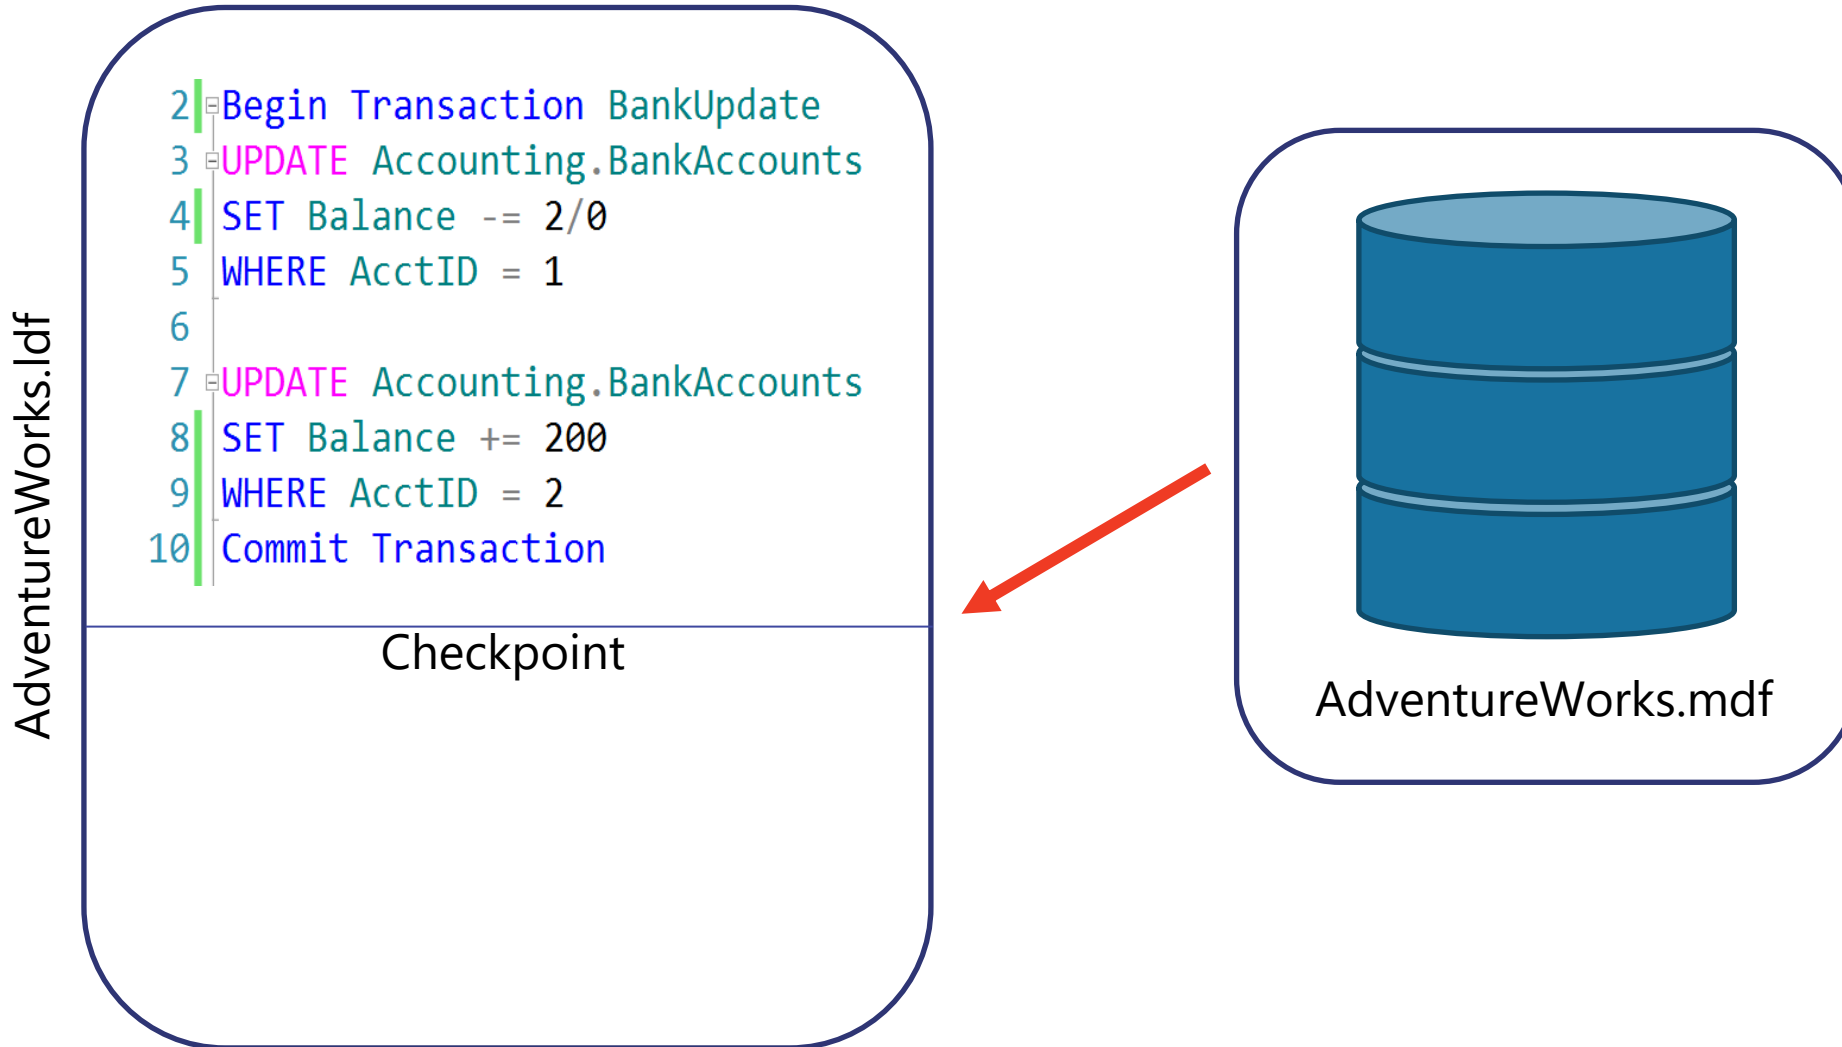
\$200

# Auto-Commit Transactions without Error Handling



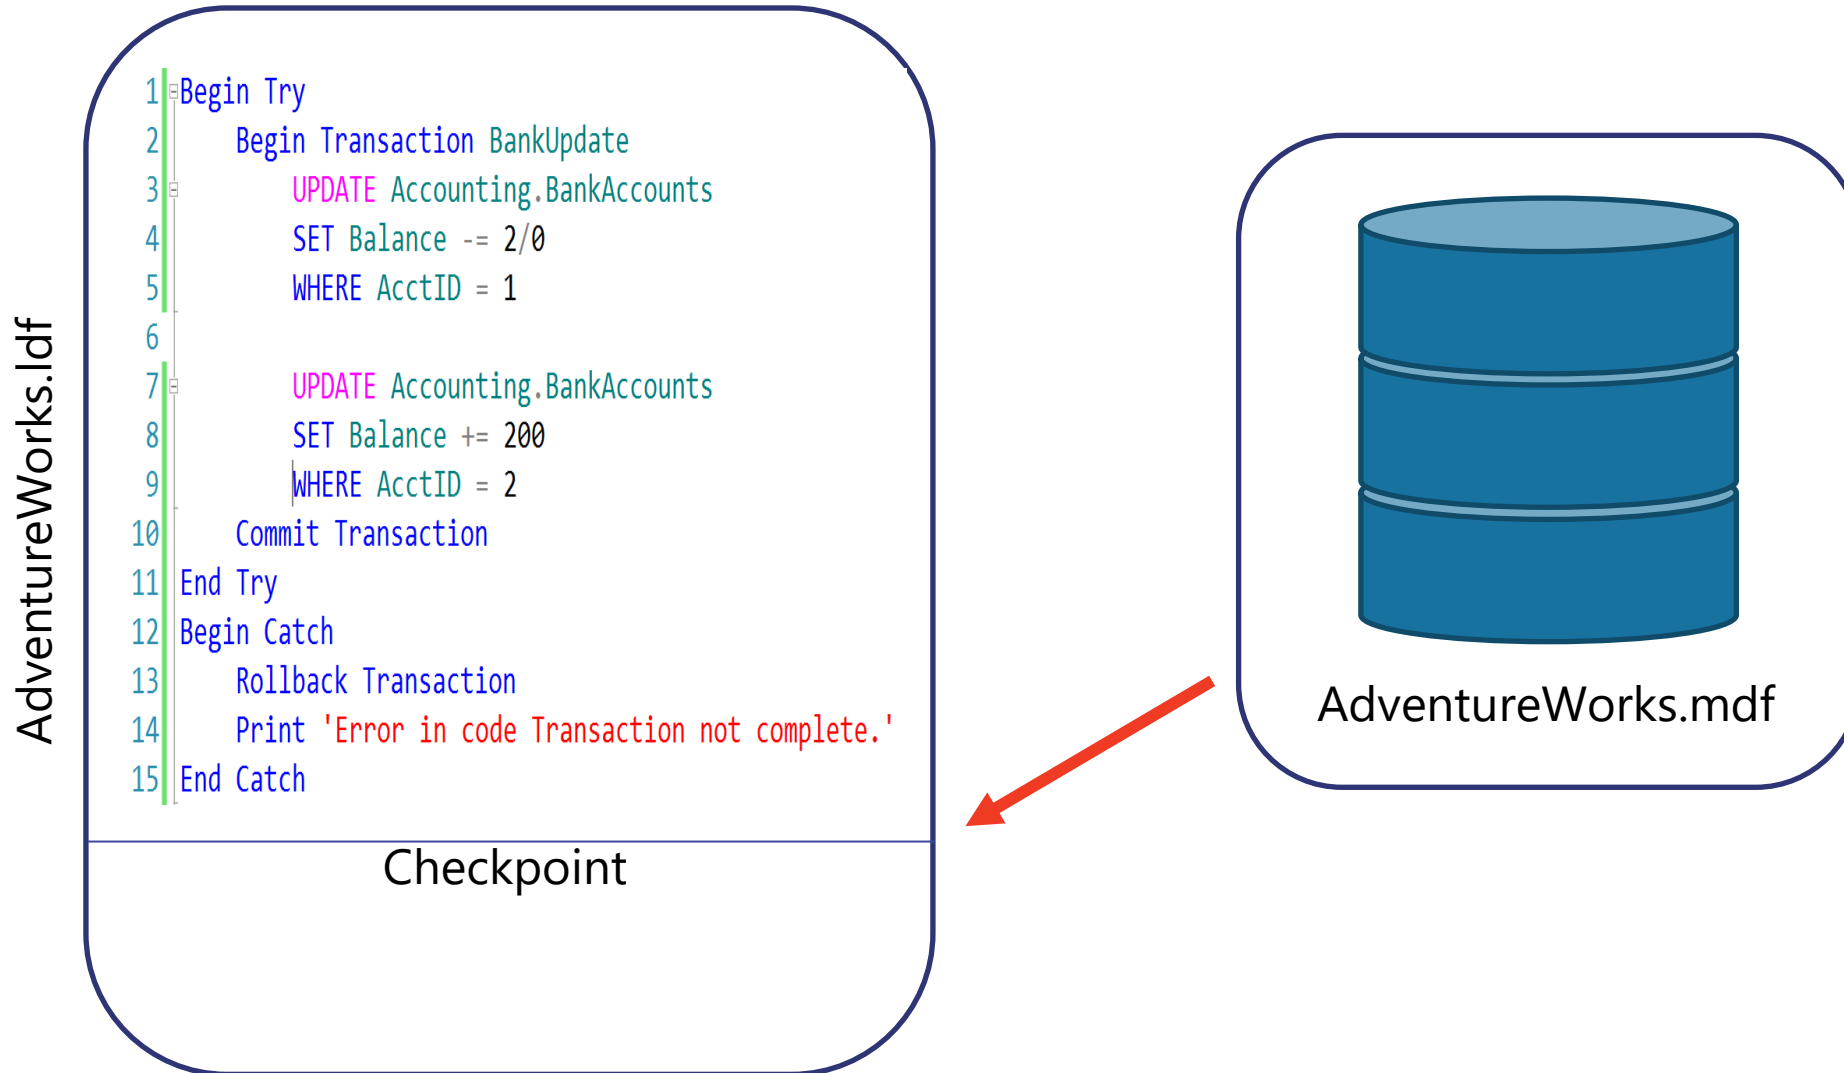
John, don't forget to demonstrate  
SET XACT\_ABORT ON

# Explicit Transactions without Error Handling





# Explicit Transactions with Error Handling



# Transactions must pass the ACID test

**Atomicity – All or Nothing**

**Consistent – Only valid data**

**Isolated – No interference**

**Durable – Data is recoverable**

# Working with Transactions

```
CREATE SCHEMA Accounting Authorization dbo
```

```
CREATE TABLE BankAccounts
```

```
(AcctID int IDENTITY,
```

```
AcctName char(15),
```

```
Balance money,
```


```
ModifiedDate date)
```

```
INSERT INTO Accounting.BankAccounts
```

```
VALUES('John', 500, GETDATE())
```

```
INSERT INTO Accounting.BankAccounts
```

```
VALUE('Jane', 750, GETDATE())
```



Messages

Msg 156, Level 15, State 1, Line 8

Incorrect syntax near the keyword 'INSERT'.

Msg 102, Level 15, State 1, Line 11

Incorrect syntax near 'VALUE'.

# Creating Stored Procedures


```
≡ ALTER PROCEDURE spAccountTransfer
  (@Amount smallmoney, @a1 tinyint, @a2 tinyint)
AS
  SET NOCOUNT ON

≡ UPDATE Accounting.BankAccounts
  SET Balance -= @Amount
  WHERE AcctID = @a1

≡ UPDATE Accounting.BankAccounts
  SET Balance += @Amount
  WHERE AcctID = @a2

PRINT 'Transfer Complete'
GO
```


# What is a Lock?



# Transaction Isolation Levels

Isolation Level	Dirty Read	Lost Update	Nonrepeatable Read	Phantoms
Read uncommitted	Yes	Yes	Yes	Yes
Read committed (default)	No	Yes	Yes	Yes
Repeatable read	No	No	No	Yes
Serializable	No	No	No	No
Snapshot	No	No	No	No

# Isolation Levels





# Lost Updates

```
1 -- SQL Server Concurrency
2 -- Lost Update - Session 1
3 USE TSQLE2012
4 GO
5 DECLARE @OldBalance int, @NewBalance int
6 BEGIN TRAN
7     SELECT @OldBalance = Balance
8     FROM Accounting.BankAccounts
9     WHERE AcctID = 1
10    SET @NewBalance = @OldBalance - 300
11    WAITFOR DELAY '00:00:30:000'
12    UPDATE Accounting.BankAccounts
13    SET Balance = @NewBalance
14    WHERE AcctID = 1
15
16    SELECT @OldBalance AS OldBalance,
17    AcctID, AcctName, Balance
18    FROM Accounting.BankAccounts
19    WHERE AcctID = 1
20 COMMIT TRAN
```

OldBalance	AcctID	AcctName	Balance
500	1	John	200.00

```
1 -- SQL Server Concurrency
2 -- Lost Update - Session 2
3 USE TSQLE2012
4 GO
5 DECLARE @OldBalance int, @NewBalance int
6 BEGIN TRAN
7     SELECT @OldBalance = Balance
8     FROM Accounting.BankAccounts
9     WHERE AcctID = 1
10    SET @NewBalance = @OldBalance - 400
11
12    UPDATE Accounting.BankAccounts
13    SET Balance = @NewBalance
14    WHERE AcctID = 1
15
16    SELECT @OldBalance AS OldBalance,
17    AcctID, AcctName, Balance
18    FROM Accounting.BankAccounts
19    WHERE AcctID = 1
20 COMMIT TRAN
```

OldBalance	AcctID	AcctName	Balance
500	1	John	100.00

# Uncommitted dependency (dirty read)

```
-- SQL Server Concurrency
-- Dirty Read - Session 1
USE TSQ2012
GO
SET TRANSACTION ISOLATION LEVEL
READ UNCOMMITTED
BEGIN TRAN
    UPDATE Accounting.BankAccounts
    SET Balance -= 300
    WHERE AcctID = 1
    WAITFOR DELAY '00:00:10:000'
    ROLLBACK TRAN
    SELECT AcctID, AcctName, Balance
    FROM Accounting.BankAccounts
    WHERE AcctID = 1
```

Clean Read

AcctID	AcctName	Balance	ModifiedDate
1	John	500.00	2013-02-16

```
--SQL Server Concurrency
--Dirty Read - Session 2
USE TSQ2012
SET TRANSACTION ISOLATION LEVEL
READ UNCOMMITTED
SELECT * FROM Accounting.BankAccounts
WHERE AcctID = 1
```

Dirty Read

AcctID	AcctName	Balance	ModifiedDate
1	John	200.00	2015-12-12

# Inconsistent analysis (non-repeatable read)

```
1 --SQL Server Concurrency
2 --Repeatable Read - Session 1
3 USE TSQL2012
4 SET TRANSACTION ISOLATION LEVEL
5 READ COMMITTED --REPEATABLE READ
6 BEGIN TRAN
7     SELECT AcctID, ModifiedDate
8     FROM Accounting.BankAccounts
9     WAITFOR DELAY '00:00:30:000'
10    SELECT AcctID, ModifiedDate
11    FROM Accounting.BankAccounts
12 COMMIT TRAN
```

```
1 --SQL Server Concurrency
2 --Repeatable Read - Session 2
3 USE TSQL2012
4 BEGIN TRAN
5     UPDATE Accounting.BankAccounts
6     SET ModifiedDate = '01/05/2013'
7 COMMIT TRAN
```

**READ COMMITTED**

AcctID	ModifiedDate
1	2015-12-12
2	2015-12-12

AcctID	ModifiedDate
1	2013-01-05
2	2013-01-05

**REPEATABLE READ**

AcctID	ModifiedDate
1	2015-12-12
2	2015-12-12

AcctID	ModifiedDate
1	2015-12-12
2	2015-12-12

# Phantom Reads

```
--SQL Server Concurrency
--Phantom Read - Session 1
USE TSQL2012
SET TRANSACTION ISOLATION LEVEL
READ COMMITTED
BEGIN TRAN
    SELECT AcctID, AcctName,
           Balance, ModifiedDate
    FROM Accounting.BankAccounts
    WAITFOR DELAY '00:00:10:000'
    SELECT AcctID, AcctName,
           Balance, ModifiedDate
    FROM Accounting.BankAccounts
COMMIT TRAN

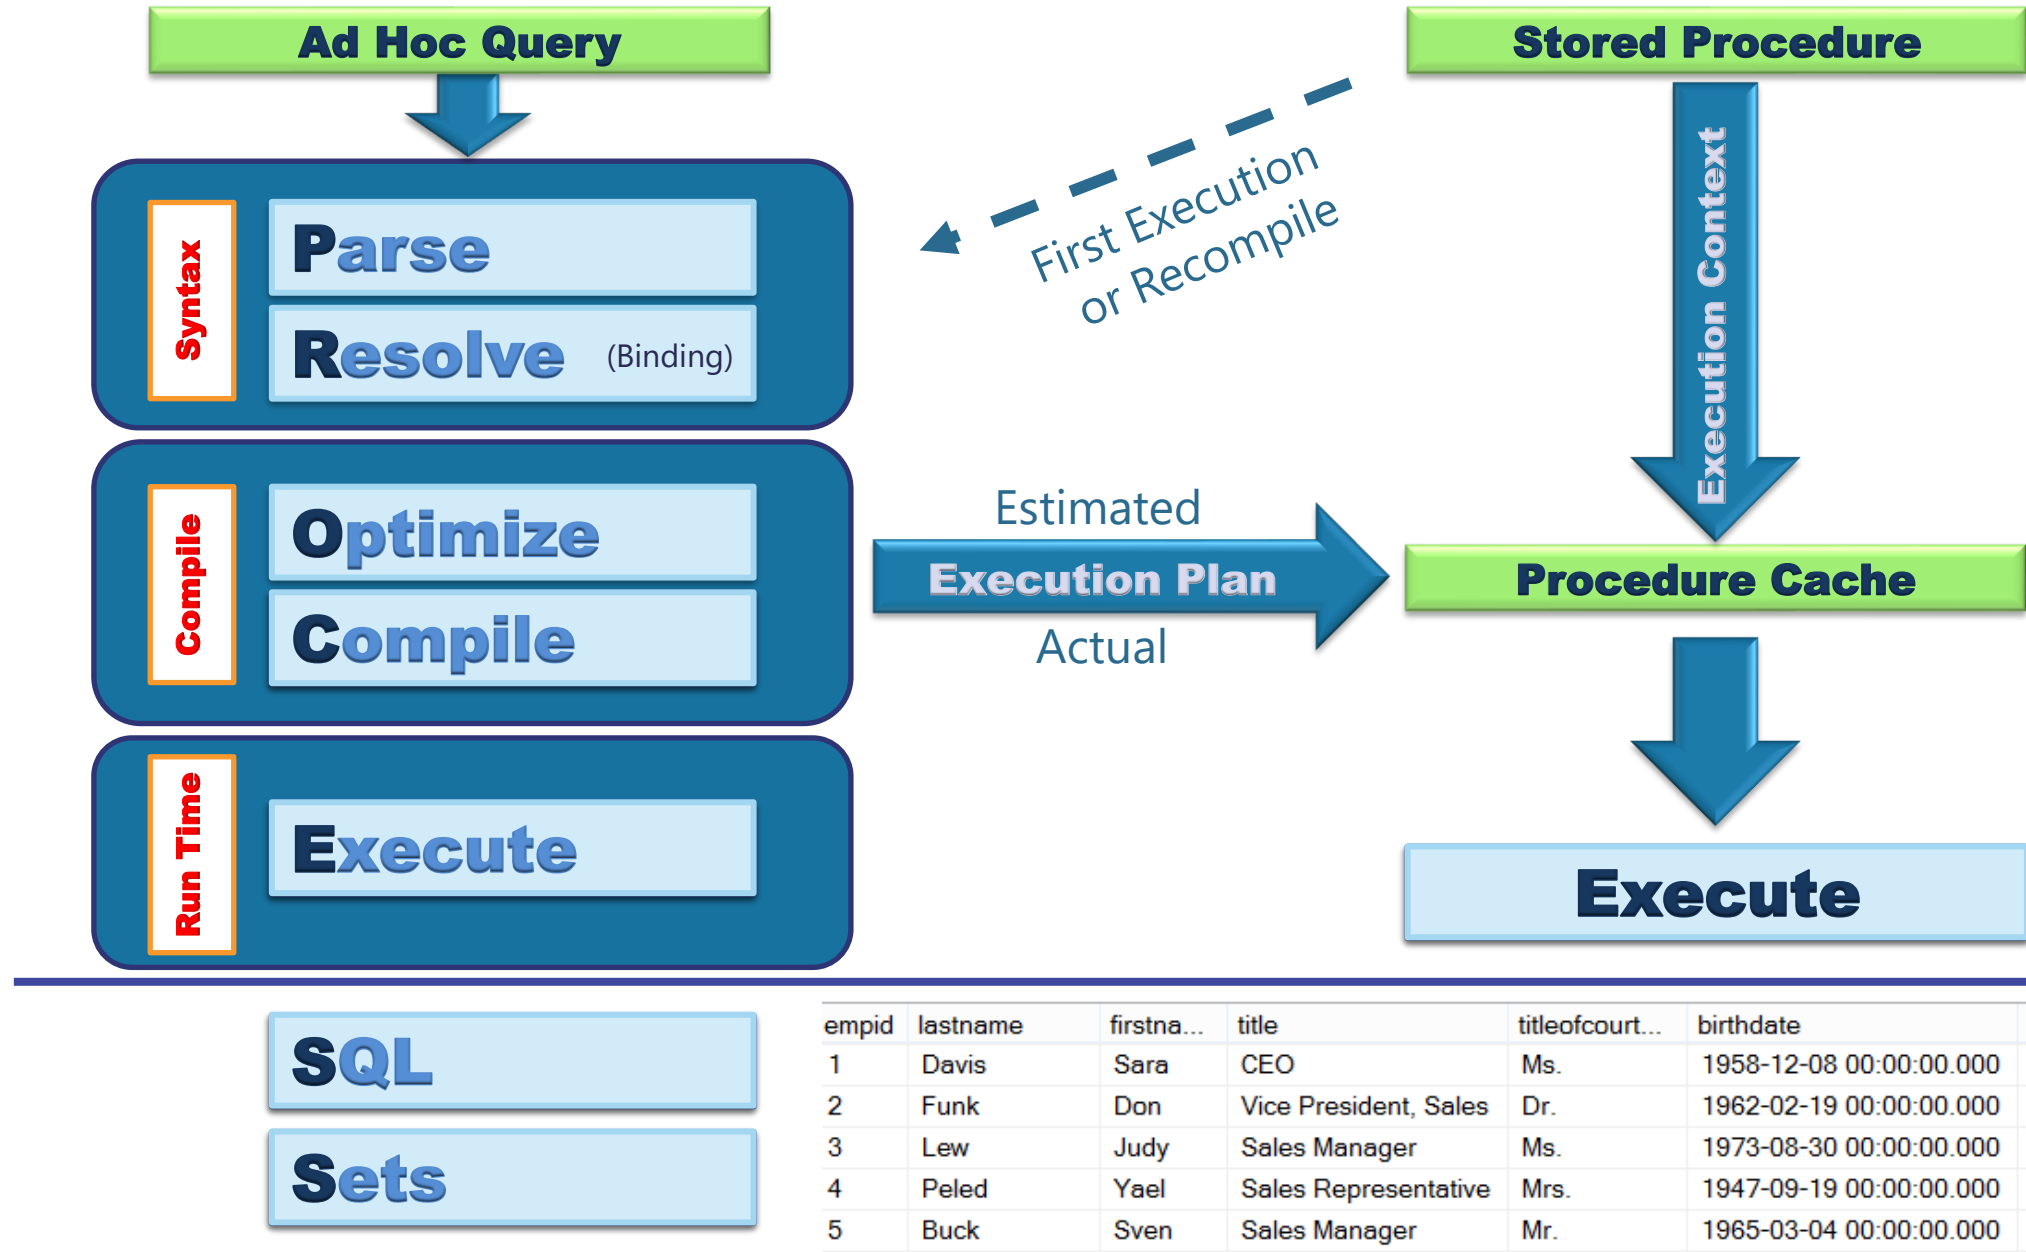
--Phantom Read - Session 2
USE TSQL2012
BEGIN TRAN
    DELETE FROM Accounting.BankAccounts
    WHERE AcctID IN(3,5,6)
COMMIT TRAN
```

Missing records

AcctID	AcctName	Balance	ModifiedDate
1	John	500.00	2016-01-02
2	Armando	750.00	2016-01-02
3	Kelli	1250.00	2016-01-02
4	Jessica	1005.00	2016-01-02
5	Maddison	745.00	2016-01-02
6	Alicen	555.00	2016-01-02
7	Molly	790.00	2016-01-02
8	Amy	650.00	2016-01-02

AcctID	AcctName	Balance	ModifiedDate
1	John	500.00	2016-01-02
2	Armando	750.00	2016-01-02
4	Jessica	1005.00	2016-01-02
7	Molly	790.00	2016-01-02
8	Amy	650.00	2016-01-02
9	Logan	1050.00	2016-01-02

# How Queries are Processed





# What is an Execution Plan?

```
SELECT SOH.SalesOrderID, SOH.CustomerID,
       OrderQty, UnitPrice, P.Name
FROM SalesLT.SalesOrderDetail
JOIN SalesLT.SalesOrderHeader
ON SOH.SalesOrderID =
JOIN SalesLT.Product
```

100 %

Messages Execution plan

Query 1: Query cost (relative to the batch): 100%

SELECT SOH.SalesOrderID, SOH.CustomerID, OrderQty,

Query executed successfully.

ready

Clustered Index Seek (Clustered)	
Scanning a particular range of rows from a clustered index.	
Physical Operation	Clustered Index Seek
Logical Operation	Clustered Index Seek
Estimated Execution Mode	Row
Storage	RowStore
Estimated Operator Cost	0.0243044 (37%)
Estimated I/O Cost	0.003125
Estimated Subtree Cost	0.0243044
Estimated CPU Cost	0.0001756
Estimated Number of Executions	32
Estimated Number of Rows	16,9375
Estimated Number of Rows to be Read	16,9375
Estimated Row Size	21 B
Ordered	True
Node ID	4

**Object**  
[AdventureWorksLT].[SalesLT].[SalesOrderDetail].

**Output List**  
[AdventureWorksLT].[SalesLT].[SalesOrderDetail].OrderQty,  
[AdventureWorksLT].[SalesLT].[SalesOrderDetail].ProductID,  
[AdventureWorksLT].[SalesLT].[SalesOrderDetail].UnitPrice

**Seek Predicates**  
Seek Keys[1]: Prefix: [AdventureWorksLT].[SalesLT].  
[SalesOrderDetail].SalesOrderID = Scalar Operator  
([AdventureWorksLT].[SalesLT].[SalesOrderHeader].  
[SalesOrderID] as [SOH].[SalesOrderID])

# Execution Plan Table Operators

Data stored in a Heap is not stored in any order and normally does not have a Primary Key.



Table Scan  
[BankAccounts]  
Cost: 100 %

Clustered Index data is stored in sorted order by the Clustering key. In many cases, this is the same value as the Primary Key.



Clustered Index Scan (Cluste...  
[BankAccounts].[pk\_acctID]  
Cost: 100 %

Using a WHERE statement on an Index could possibly have the Execution Plan seek the Index instead of scan.




Clustered Index Seek (Cluste...  
[BankAccounts].[pk\_acctID]  
Cost: 100 %

# Execution Plan Join Operators (Code)

```
SELECT SOH.SalesOrderID, SOH.CustomerID,  
       OrderQty, UnitPrice, P.Name  
FROM Sales.SalesOrderHeader AS SOH  
     JOIN Sales.SalesOrderDetail AS SOD  
         ON SOH.SalesOrderID = SOD.SalesOrderID  
     JOIN Production.Product AS P  
         ON P.ProductID = SOD.ProductID
```



# Execution Plan Join Operators (Plan)



# Execution Plan Join Operators

A Merge Join is useful if both table inputs are in the same sorted order on the same value.



Merge Join  
(Inner Join)  
Cost: 39 %

A Hash Match is used when the tables being joined are not in the same sorted order.



Hash Match  
(Inner Join)  
Cost: 47 %

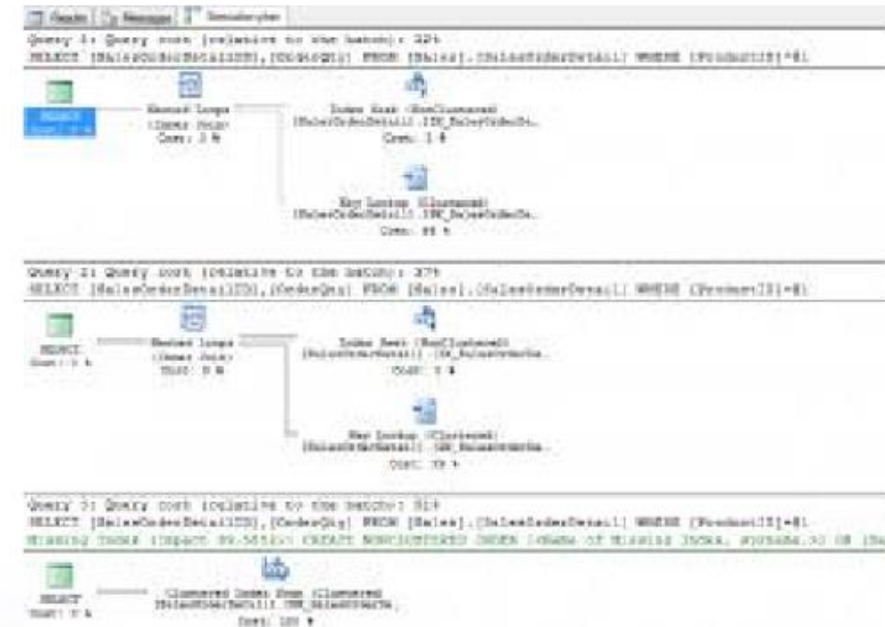
A Nested Loop is used when a small (outer) table is used to lookup a value in a larger (inner) table.



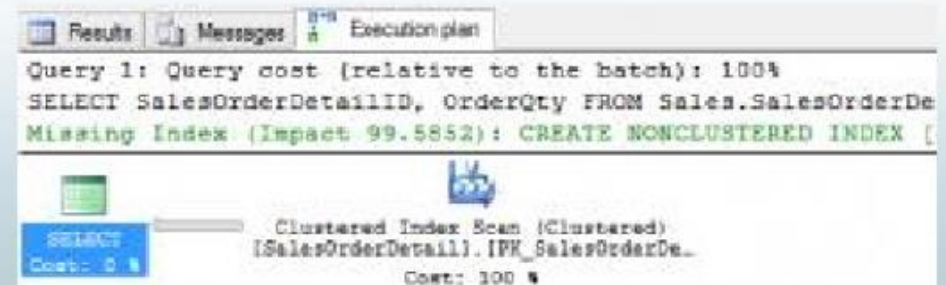
Nested Loops  
(Inner Join)  
Cost: 3 %

# Parameter Sniffing

```
1 SELECT SalesOrderDetailID, OrderQty
2 FROM Sales.SalesOrderDetail
3 WHERE ProductID = 897
4
5 SELECT SalesOrderDetailID, OrderQty
6 FROM Sales.SalesOrderDetail
7 WHERE ProductID = 945
8
9 SELECT SalesOrderDetailID, OrderQty
10 FROM Sales.SalesOrderDetail
11 WHERE ProductID = 870
```



```
1 CREATE PROCEDURE Get_OrderQuantity
2 (@ProductID int)
3 AS
4 SELECT SalesOrderDetailID, OrderQty
5 FROM Sales.SalesOrderDetail
6 WHERE ProductID = @ProductID
```



# Cardinality Estimator and Statistics

Statistics Properties - IX\_SalesOrderDetail\_ProductID

Select a page

- General
- Details
- Filter

Script Help

Table Name: SalesLT.SalesOrderDetail

Statistics Name: IX\_SalesOrderDetail\_ProductID

Statistics for INDEX 'IX\_SalesOrderDetail\_ProductID'.

Name	Updated
IX_SalesOrderDetail_ProductID	Aug 20 2019 1:09PM

All Density	Average Length
0.007042253	4
0.001845018	8
0.001845018	12

Pr

ProductID,

Histogram Steps	RANGE_HI_KEY	RANGE_ROWS
707		0
708		0
711		0
712		0
714		0
715		0
716		0
718		2
722		0
738		0
739		0
742		0
743		0
747		0
748		0

Ready

OK Cancel

